

W01

3.1 Problem-Solving Agents

Formulating problems

Search problems and solutions

3.2 Example problems

Standard problem

Grid world
Vacuum world
Sokoban puzzle
Sliding-tile puzzle
Knuth's 4

Real world problem

Route-finding problem
TSP
Touring problems
VLSI problem
Robot navigation
Automatic assembly sequencing

3.3 Search Algorithms

Best-first search

Search data structures

Priority queue
FIFO: queue
LIFO: stack

Redundant paths

Graph search: check redundant
Check limited ancestor
Assembly problem
Tree search: does not check redundant
BFS
Remember all previously reached states
Rely on other mechanism to deal with long cycles

Measuring problem-solving performance

Completeness
Cost optimality

IVI
IEI

Time and Space

d: depth
m: max actions in any path
b: branching factor

3.4 Uniformed Search Strategies

BFS
Dijkstra or uniform-cost search
Backtracking search
Depth-limited and iterative deepening search
Bidirectional search

Comparing

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ¹	Yes ^{1,2}	No	No	Yes ¹	Yes ^{1,4}
Optimal cost?	Yes ³	Yes	No	No	Yes ³	Yes ^{3,4}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b^l)$	$O(bd)$	$O(b^{d/2})$

Informed(Heuristic) Search Strategies

Greedy best first search
Admissible heuristic: never overestimates
Prunes nodes that are not necessary for optimal
A* Search
Search contours
Satisficing search
Weighted A*
Beam search
Memory-bounded search
IDA*
RBFS
SMA*
Bidirectional A*
Bidirectional heuristic search
Effective depth
The effect of heuristic accuracy on performance

Heuristic Functions

Generating heuristics from relaxed problems
Generating heuristics from sub-problems: pattern databases
Pattern database
Disjoint pattern databases
Difference to DP problem???

Generating heuristics with landmarks
Learning to search better
Learning heuristics from experience

State Space Search

All of AI is search

4 key elements

States node
Actions
Transition model arc
Cost (action cost function)

Solution may not unique

Problem: Episodic, single agent, fully observable, deterministic, static, discrete, known

Generate the state when we need

size of the state space graph

State: dictionary, list
Transition: Function, table, map
Location as values

Frame the problem

Problem to resolve: **planning**, routing, design, sequencing, puzzles, and **games**

Uninformed search

Agent can not estimate how far it is from the goal

Complete N
Optimal N

Space: $O(bm)$
b: branching factor
d: depth of shallowest solution
m: max depth of the resulting search tree

Time (operations, not CPU): $O(b^m)$

Improvement: iterative deepening => why not BFS?

Complete Y
Optimal Y (uniform cost)

Space: $O(b^m)$

Time: $O(b^m)$

Declarative programming

Function, logic, DB approach
expresses the logic of a computation without describing its control flow
What to do: Red? Blue: red
Imperative programming: How to do: if RED, blue. If Blue, red

Informed search

The cost model may give us important information about good solution and where to look for them

UCS, Uniform cost search (Priority queue)

Cost: as our key. Examine the node on the frontier with lowest total path cost so far
Successor function with cost-calculated annotation
BFS++

Greedy(Priority queue)

Best first search
A simplest variant: pick the cheapest action **at the current**.
Cost estimates come from heuristic functions
Might faster than UCS
Doesn't always work
Might deadlock itself

Estimate current node to the goal
Lecture: use current action cost as an estimate
c: the shortest cost from previous node, through action node, to the goal
Admissible heuristic: $h(n) \leq c$ The best $h(n)$ is the c, better is close to c

Complete (finite state space), not optimal.
Time/space: $O(|V|)$ or some cases $O(bm)$

A* search (Priority queue)

UCS+Greedy
Extremely useful for path finding
 $f(n) = g(n) + h(n)$
g(n): cost so far, from UCS
h(n): heuristic from greedy
The estimated cost of getting from the current node to the goal
Optimal (may or may not) and complete
If there is a path to the goal
Depends on heuristic function