



W13

Perceptron (13A)

- Connectionist: Simulate Brian instead of high level symbols and logic
- Perceptron
  - Typical use 1 and -1
  - Threshold function

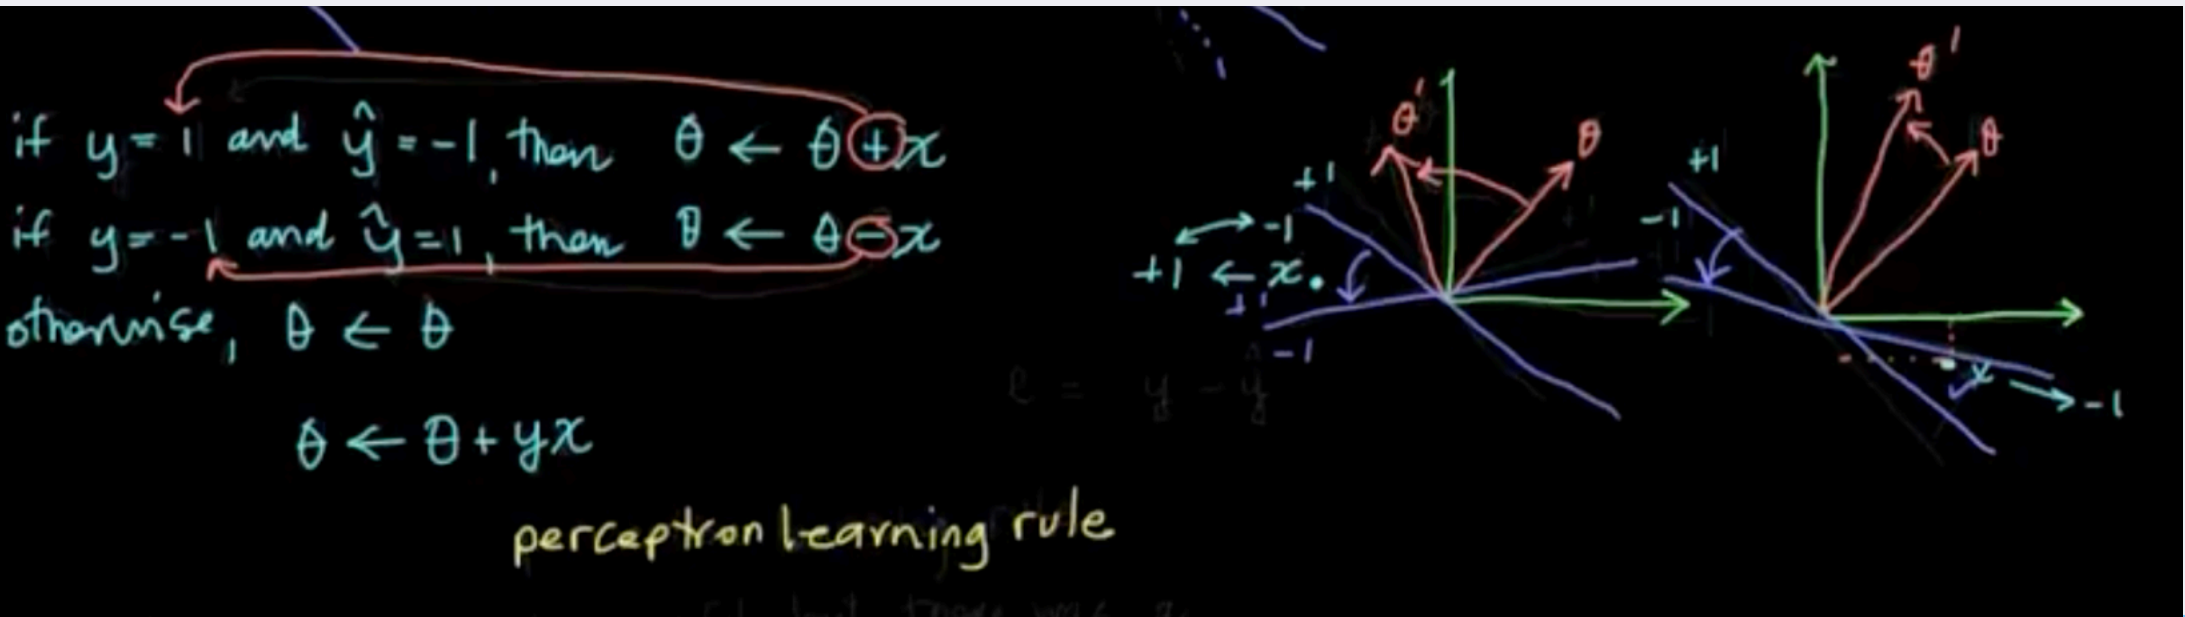


The decision boundary is orthogonal to the weights



Perceptron Learning Rule (13B)

Basic vector arithmetic



Perceptron learning rule

A single line can't represent a non-linearly separable concept

Multi-Layer Perceptrons (13C)

Progress

MLP

Adaline

- NN version of regression
- Stochastic gradient descent
- Transfer function is linear, rather than threshold
- Noise cancellation

Clamps to 0-1 range (or -1 to 1), for nice derivative

Stack them and run in parallel

- Input, hidden, output
- 1 hidden layer, can approximate any continuous function
- 2 hidden layers, can approximate any function

Kernel trick

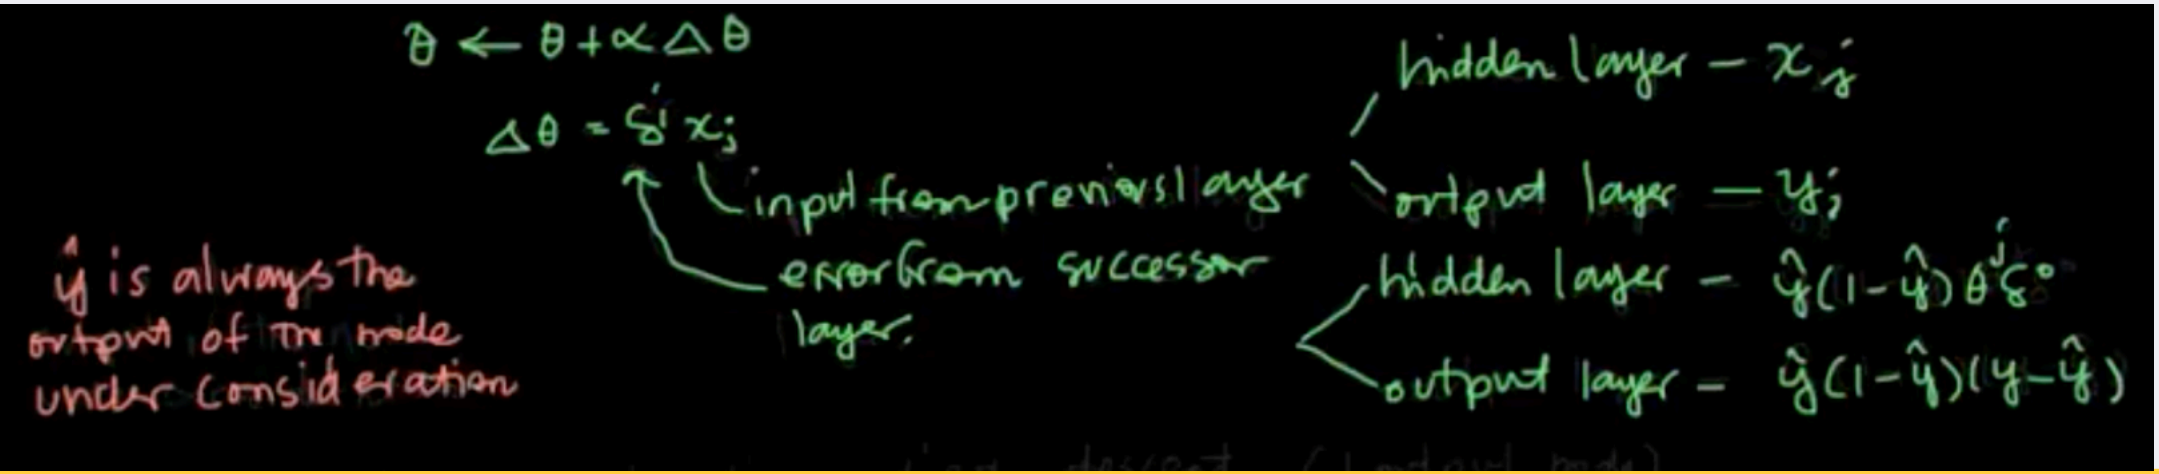
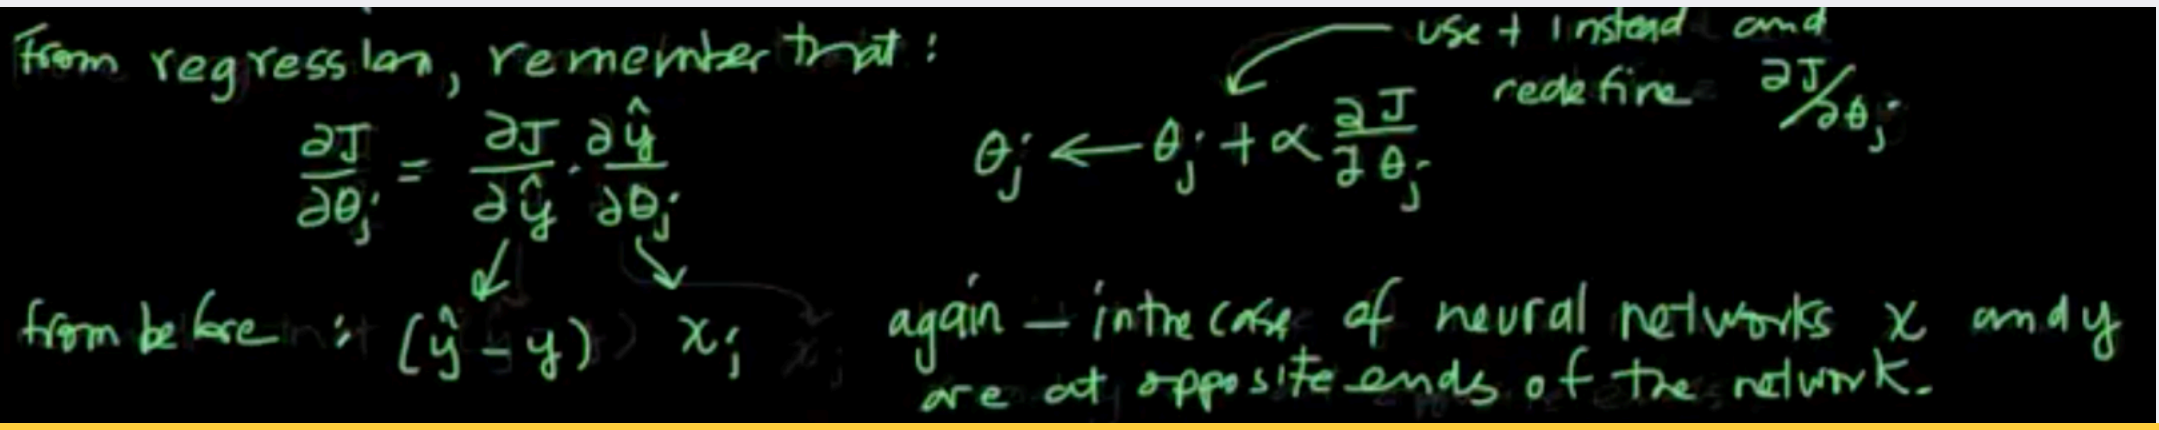
Based on a dot product, voted perceptron

Feed forward NN

Main problem: training data only have input, not hidden

Backpropagation Algorithm (13D)

Back propagation algorithm



Signal go forward, error go backward

Multiply output

$$\delta^h = \hat{y} (1 - \hat{y}) \sum \theta_i \delta_i^o$$

Stop when changing of error is small

Tips for Using Multi-Layer Perceptrons (13E)

Initialize all weights to small random values, 0-1

Scale all input

Mean scaling

Stuck in local minima

Random restarts

Optimize techniques

- Genetic algorithm
- Conjugate gradient descent
- combinations

Pick the number of hidden layers + number of nodes in each layer

- Start simple
- Use cross validation

Very easy overfitting

Prone, use regulation

Multi-class problem

Binary encoding

00, 01, 10, 11

Categorical input